

Introduction to Machine Learning, Fall 2023

Homework 1

(Due Thursday, Oct. 26 at 11:59pm (CST))

October 25, 2023

1. [10 points] [Math review] Suppose $\{\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n\}$ are random samples from a random variable \mathbf{X} :

- (a) Prove that the covariance of \mathbf{X} is a semi positive definite matrix. [3 points]
- (b) Assuming $\mathbf{X} \sim \mathcal{N}(\mu, \Sigma)$ which is a multivariate normal distribution, derive the log-likelihood $l(\mu, \Sigma)$ and MLE of μ [4 points]
- (c) Suppose $\hat{\theta}$ is an unbiased estimator of θ and $\text{Var}(\hat{\theta}) > 0$. Prove that $(\hat{\theta})^2$ is not an unbiased estimator of θ^2 . [3 points]

(a) $\Sigma = E[(\mathbf{X} - \mu)(\mathbf{X} - \mu)^T]$

So $\forall y \in \mathbb{R}^n$,

$$y^T \Sigma y = E[y^T (\mathbf{X} - \mu)(\mathbf{X} - \mu)^T y] = E[(\mathbf{X} - \mu)^T y]^T ((\mathbf{X} - \mu)^T y) = E(\|(\mathbf{X} - \mu)^T y\|_2^2).$$

since $\|(\mathbf{X} - \mu)^T y\|_2^2 \geq 0$, so $E(\|(\mathbf{X} - \mu)^T y\|_2^2) \geq 0$

so $\forall y \in \mathbb{R}^n, y^T \Sigma y \geq 0$

So Σ is a semi positive definite matrix.

(b) From what we have learned in class, the PDF of the multivariate normal distribution $\mathbf{X}_i \sim N(\mu, \Sigma)$ is that $Pr(\mathbf{X}_i; \mu, \Sigma) = \frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu))$, suppose that the dimension of \mathbf{X} is p .

Since the sampling are independent, so the likelihood function is:

$$Pr(\mathbf{X}_1, \dots, \mathbf{X}_n; \mu, \Sigma) = \prod_{i=1}^n Pr(\mathbf{X}_i; \mu, \Sigma)$$

Then the log-likelihood function is:

$$\begin{aligned} l(\mu, \Sigma) &= \sum_{i=1}^n \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu))\right) \\ &= n \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}}\right) - \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu) \end{aligned}$$

And the MLE of μ is:

$$\hat{\mu} = \underset{\mu}{\operatorname{argmax}} Pr(\mathbf{X}_1, \dots, \mathbf{X}_n; \mu, \Sigma) = \underset{\mu}{\operatorname{argmax}} l(\mu, \Sigma)$$

Since $l(\mu, \Sigma)$ is a concave function, so we can get the optimal solution by setting the derivative of $l(\mu, \Sigma)$ to 0.

$$\begin{aligned} \frac{\partial l(\mu, \Sigma)}{\partial \mu} &= \frac{\partial b \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}}\right)}{\partial \mu} - \frac{\partial \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu)}{\partial \mu} \\ &= \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu)}{\partial \mu} \end{aligned}$$

Since Σ is the covariance matrix, so Σ is a symmetric matrix, i.e. $\Sigma = \Sigma^T$.

So $(\Sigma^{-1})^T = (\Sigma^T)^{-1} = \Sigma^{-1}$.

i.e. Σ is also a symmetric matrix.

$$\begin{aligned} \text{So } \frac{\partial l(\mu, \Sigma)}{\partial \mu} &= \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu)}{\partial \mu} \\ &= \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma (\mathbf{X}_i - \mu)}{\partial(\mathbf{X}_i - \mu)} \frac{\partial(\mathbf{X}_i - \mu)}{\partial \mu} \end{aligned}$$

$$\begin{aligned}
&= \sum_{i=1}^n -\frac{1}{2}(2\mathbf{\Sigma}(\mathbf{X} - \mu))(-1) \\
&= \sum_{i=1}^n \mathbf{\Sigma}(\mathbf{X}_i - \mu) \\
&= \mathbf{\Sigma}(\sum_{i=1}^n \mathbf{X}_i - n\mu)
\end{aligned}$$

$$\text{So } \frac{\partial l(\mu, \mathbf{\Sigma})}{\partial \mu} = 0 \Rightarrow \mathbf{\Sigma}(\sum_{i=1}^n \mathbf{X}_i - n\mu) = 0 \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$$

$$\text{So above all, the log-likelihood function is } l(\mu, \mathbf{\Sigma}) = n \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\mathbf{\Sigma}|^{\frac{1}{2}}}\right) - \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \mathbf{\Sigma}(\mathbf{X}_i - \mu)$$

$$\text{And the MLE of } \mu \text{ is } \hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$$

(c)

2. [10 points] Consider real-valued variables X and Y , in which Y is generated conditional on X according to

$$Y = aX + b + \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2).$$

Here ϵ is an independent variable, called a noise term, which is drawn from a Gaussian distribution with mean 0, and variance σ^2 . This is a single variable linear regression model, where a is the only weight parameter and b denotes the intercept. The conditional probability of Y has a distribution $p(Y|X, a, b) \sim \mathcal{N}(aX + b, \sigma^2)$, so it can be written as:

$$p(Y|X, a, b) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(Y - aX - b)^2\right).$$

- (a) Assume we have a training dataset of n i.i.d. pairs (x_i, y_i) , $i = 1, 2, \dots, n$, and the likelihood function is defined by $L(a, b) = \prod_{i=1}^n p(y_i|x_i, a, b)$. Please write the Maximum Likelihood Estimation (MLE) problem for estimating a and b . [3 points]
 - (b) Estimate the optimal solution of a and b by solving the MLE problem in (a). [4 points]
 - (c) Based on the result in (b), argue that the learned linear model $f(X) = aX + b$, always passes through the point (\bar{x}, \bar{y}) , where $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$ and $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$ denote the sample means. [3 points]
- (a)
- (b)
- (c)

3. [10 points] [Regression and Classification]

- (a) When we talk about linear regression, what does ‘linear’ regard to? [2 points]
- (b) Assume that there are n given training examples $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$, where each input data point x_i has m real valued features. When $m > n$, the linear regression model is equivalent to solving an under-determined system of linear equations $\mathbf{y} = \mathbf{X}\beta$. One popular way to estimate β is to consider the so-called ridge regression:

$$\underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2$$

for some $\lambda > 0$. This is also known as Tikhonov regularization.

Show that the optimal solution β_* to the above optimization problem is given by

$$\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Hint: You need to prove that given $\lambda > 0$, $\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}$ is invertible. [5 points]

- (c) Is the given data set linear separable? If yes, construct a linear hypothesis function to separate the given data set. If no, explain the reason. [3 points]

Data	(1,3)	(4,4)	(3,-6)	(-2,1)	(-3,5)	(-6,-4)
Label	+1	-1	-1	+1	-1	-1

(a) Linear

(b) As we have learned in linear algebra, we know that the matrix $X^T X$ must be similar and diagonalizable. i.e. there must exist a matrix P and a diagonal matrix Λ such that $X^T X = P \Lambda P^{-1}$.

Also $\forall x \in \mathbb{R}^n$, we have $x^T (\mathbf{X}^T \mathbf{X}) x = (\mathbf{X}x)^T (\mathbf{X}x) = \|\mathbf{X}x\|_2^2 \geq 0$.

So $X^T X$ is positive semi-definite.

So all eigenvalues of $X^T X$ are non-negative.

i.e. the diagonal matrix Λ 's elements are all positive.

And since $\lambda > 0$, so λI 's all elements are all also non-negative, and λI is also a diagonal matrix.

So $X^T X + \lambda I = P \Lambda P^{-1} + \lambda P I P^{-1} = P(\Lambda + \lambda I) P^{-1}$.

Since $\Lambda, \lambda I$ are all diagonal matrix, so $\Lambda + \lambda I$ is also a diagonal matrix.

And all elements in $\Lambda + \lambda I$ are all positive, this is because in Λ , elements are non-negative, in λI , all elements are positive. So $\Lambda + \lambda I$ is positive defined.

Since $X^T X + \lambda I = P(\Lambda + \lambda I) P^{-1}$, from the knowledge of similarity and diagonalizable, we could know that $X^T X + \lambda I$ is also positive defined.

So $X^T X + \lambda I$ is invertible.

And let $f(\beta) = \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2 = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) + \lambda \beta^T \beta = \mathbf{y}^T \mathbf{y} - \beta^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \beta + \beta^T \mathbf{X}^T \mathbf{X} \beta + \lambda \beta^T \beta$

Since $f(\beta)$ is convex, so we just need to set the derivative of $f(\beta)$ to 0 to get the optimal solution.

$$\frac{\partial f(\beta)}{\partial \beta} = 2(-\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \beta + \lambda \beta)$$

$$\frac{\partial f(\beta)}{\partial \beta} = 0 \Rightarrow (\mathbf{X}^T \mathbf{X} + \lambda I) \beta = \mathbf{X}^T \mathbf{y} \Rightarrow \beta = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

Since we have proved that $X^T X + \lambda I$ is invertible, so $(\mathbf{X}^T \mathbf{X} + \lambda I)^{-1}$ exists. So $\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$

So above all, the optimal solution $\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$.

(c)