

# Introduction to Machine Learning, Fall 2023

## Homework 1

(Due Thursday, Oct. 26 at 11:59pm (CST))

October 25, 2023

1. [10 points] [Math review] Suppose  $\{\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_n\}$  are random samples from a random variable  $\mathbf{X}$ :

- (a) Prove that the covariance of  $\mathbf{X}$  is a semi positive definite matrix. [3 points]
- (b) Assuming  $\mathbf{X} \sim \mathcal{N}(\mu, \Sigma)$  which is a multivariate normal distribution, derive the the log-likelihood  $l(\mu, \Sigma)$  and MLE of  $\mu$  [4 points]
- (c) Suppose  $\hat{\theta}$  is an unbiased estimator of  $\theta$  and  $\text{Var}(\hat{\theta}) > 0$ . Prove that  $(\hat{\theta})^2$  is not an unbiased estimator of  $\theta^2$ . [3 points]

(a)  $\Sigma = E[(\mathbf{X} - \mu)(\mathbf{X} - \mu)^T]$

Suppose that the dimension of  $\mathbf{X}$  is  $p$ .

So  $\forall \mathbf{y} \in \mathbb{R}^p$ ,

$$\mathbf{y}^T \Sigma \mathbf{y} = \mathbf{y}^T E[(\mathbf{X} - \mu)(\mathbf{X} - \mu)^T] \mathbf{y} = E[\mathbf{y}^T (\mathbf{X} - \mu)(\mathbf{X} - \mu)^T \mathbf{y}]$$

$$= E[(\mathbf{X} - \mu)^T \mathbf{y}]^T (\mathbf{X} - \mu)^T \mathbf{y}] = E(\|\mathbf{X} - \mu\|^2 \mathbf{y}^T \mathbf{y}).$$

since  $\|\mathbf{X} - \mu\|^2 \geq 0$ , so  $E(\|\mathbf{X} - \mu\|^2 \mathbf{y}^T \mathbf{y}) \geq 0$

so  $\forall \mathbf{y} \in \mathbb{R}^p, \mathbf{y}^T \Sigma \mathbf{y} \geq 0$

So  $\Sigma$  is a semi positive definite matrix.

So above all, the covariance of  $\mathbf{X}, \Sigma$  is a semi positive definite matrix.

(b) From what we have learned in class, the PDF of the multivariate normal distribution  $\mathbf{X}_i \sim \mathcal{N}(\mu, \Sigma)$  is that  $Pr(\mathbf{X}_i; \mu, \Sigma) = \frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu))$ , suppose that the dimension of  $\mathbf{X}_i$  is  $p$ .

Since the sampling are independent, so the likelihood function is:

$$Pr(\mathbf{X}_1, \dots, \mathbf{X}_n; \mu, \Sigma) = \prod_{i=1}^n Pr(\mathbf{X}_i; \mu, \Sigma)$$

Then the log-likelihood function is:

$$l(\mu, \Sigma) = \sum_{i=1}^n \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}} \exp(-\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu))\right)$$

$$= n \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}}\right) - \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu)$$

And the MLE of  $\mu$  is:

$$\hat{\mu} = \underset{\mu}{\operatorname{argmax}} Pr(\mathbf{X}_1, \dots, \mathbf{X}_n; \mu, \Sigma) = \underset{\mu}{\operatorname{argmax}} l(\mu, \Sigma)$$

Since  $l(\mu, \Sigma)$  is a concave function, so we can get the optimal solution by setting the derivative of  $l(\mu, \Sigma)$  to 0.

$$\begin{aligned} \frac{\partial l(\mu, \Sigma)}{\partial \mu} &= \frac{\partial n \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\Sigma|^{\frac{1}{2}}}\right)}{\partial \mu} - \frac{\partial \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu)}{\partial \mu} \\ &= \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu)}{\partial \mu} \end{aligned}$$

Since  $\Sigma$  is the covariance matrix, so  $\Sigma$  is a symmetric matrix, i.e.  $\Sigma = \Sigma^T$ .

So  $(\Sigma^{-1})^T = (\Sigma^T)^{-1} = \Sigma^{-1}$ .

i.e.  $\Sigma$  is also a symmetric matrix.

$$\text{So } \frac{\partial l(\mu, \Sigma)}{\partial \mu} = \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \Sigma^{-1}(\mathbf{X}_i - \mu)}{\partial \mu}$$

$$\begin{aligned}
&= \sum_{i=1}^n -\frac{\frac{1}{2}(\mathbf{X}_i - \mu)^T \boldsymbol{\Sigma}^{-1}(\mathbf{X}_i - \mu)}{\frac{\partial(\mathbf{X}_i - \mu)}{\partial \mu}} \frac{\partial(\mathbf{X}_i - \mu)}{\partial \mu} \\
&= \sum_{i=1}^n -\frac{1}{2}(2\boldsymbol{\Sigma}^{-1}(\mathbf{X} - \mu))(-1) \\
&= \sum_{i=1}^n \boldsymbol{\Sigma}^{-1}(\mathbf{X}_i - \mu) \\
&= \boldsymbol{\Sigma}^{-1}\left(\sum_{i=1}^n \mathbf{X}_i - n\mu\right)
\end{aligned}$$

$$\text{So } \frac{\partial l(\mu, \boldsymbol{\Sigma})}{\partial \mu} = 0 \Rightarrow \boldsymbol{\Sigma}^{-1}\left(\sum_{i=1}^n \mathbf{X}_i - n\mu\right) = 0 \Rightarrow \hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$$

So above all, the log-likelihood function is  $l(\mu, \boldsymbol{\Sigma}) = n \cdot \log\left(\frac{1}{(2\pi)^{\frac{p}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}}\right) - \sum_{i=1}^n \frac{1}{2}(\mathbf{X}_i - \mu)^T \boldsymbol{\Sigma}^{-1}(\mathbf{X}_i - \mu)$

And the MLE of  $\mu$  is  $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i$

(c) Since  $\hat{\theta}$  is the unbiased estimator of  $\theta$ , so  $E(\hat{\theta}) = \theta$ .

And from the definition of variance, we could get that  $Var(\hat{\theta}) = E[(\hat{\theta} - E(\hat{\theta}))^2] = E[(\hat{\theta})^2] - [E(\hat{\theta})]^2$ .

Since  $Var(\hat{\theta}) > 0$ , so  $E[(\hat{\theta})^2] - [E(\hat{\theta})]^2 > 0$

i.e.  $E[(\hat{\theta})^2] > [E(\hat{\theta})]^2 = (\hat{\theta})^2$

So  $E[(\hat{\theta})^2] \neq (\hat{\theta})^2$

So  $(\hat{\theta})^2$  is not an unbiased estimator of  $\theta^2$ .

So above all, we have proved that if  $\hat{\theta}$  is an unbiased estimator of  $\theta$  and  $Var(\hat{\theta}) > 0$ , then  $(\hat{\theta})^2$  is not an unbiased estimator of  $\theta^2$ .

2. [10 points] Consider real-valued variables  $X$  and  $Y$ , in which  $Y$  is generated conditional on  $X$  according to

$$Y = aX + b + \epsilon, \text{ where } \epsilon \sim \mathcal{N}(0, \sigma^2).$$

Here  $\epsilon$  is an independent variable, called a noise term, which is drawn from a Gaussian distribution with mean 0, and variance  $\sigma^2$ . This is a single variable linear regression model, where  $a$  is the only weight parameter and  $b$  denotes the intercept. The conditional probability of  $Y$  has a distribution  $p(Y|X, a, b) \sim \mathcal{N}(aX + b, \sigma^2)$ , so it can be written as:

$$p(Y|X, a, b) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(Y - aX - b)^2\right).$$

- (a) Assume we have a training dataset of  $n$  i.i.d. pairs  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ , and the likelihood function is defined by  $L(a, b) = \prod_{i=1}^n p(y_i|x_i, a, b)$ . Please write the Maximum Likelihood Estimation (MLE) problem for estimating  $a$  and  $b$ . [3 points]
- (b) Estimate the optimal solution of  $a$  and  $b$  by solving the MLE problem in (a). [4 points]
- (c) Based on the result in (b), argue that the learned linear model  $f(X) = aX + b$ , always passes through the point  $(\bar{x}, \bar{y})$ , where  $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$  and  $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$  denote the sample means. [3 points]

(a) the MLE of  $a$  and  $b$  is:

$$\hat{a}, \hat{b} = \underset{a, b}{\operatorname{argmax}} L(a, b) = \underset{a, b}{\operatorname{argmax}} \prod_{i=1}^n p(y_i|x_i, a, b) = \underset{a, b}{\operatorname{argmax}} \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(y_i - ax_i - b)^2\right)$$

So above all, the MLE problem for estimating  $a$  and  $b$  is:

$$\hat{a}, \hat{b} = \underset{a, b}{\operatorname{argmax}} \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(y_i - ax_i - b)^2\right)$$

(b) Take log to the likelihood function, we could get:

$$\hat{a}, \hat{b} = \underset{a, b}{\operatorname{argmax}} \sum_{i=1}^n \log(p(y_i|x_i, a, b)) = \underset{a, b}{\operatorname{argmax}} \sum_{i=1}^n \log\left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(y_i - ax_i - b)^2\right)\right)$$

Since  $\frac{1}{\sqrt{2\pi}\sigma}$  has nothing with  $a, b$ , and  $\sigma$  is just the variance of the noise term, so  $\frac{1}{\sqrt{2\pi}\sigma}, -\frac{1}{2\sigma^2}$  are just constants.

$$\text{so } \hat{a}, \hat{b} = \underset{a, b}{\operatorname{argmax}} \sum_{i=1}^n -(y_i - ax_i - b)^2 = \underset{a, b}{\operatorname{argmin}} \sum_{i=1}^n (y_i - ax_i - b)^2$$

Since  $\sum_{i=1}^n (y_i - ax_i - b)^2$  is a convex function both for  $a$  and  $b$ , so we just need to set the derivative of

$\sum_{i=1}^n (y_i - ax_i - b)^2$  to 0 to get the optimal solution.

$$\text{Let } f(a, b) = \sum_{i=1}^n (y_i - ax_i - b)^2, \bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$$

$$\text{So } \frac{\partial f}{\partial b} = \sum_{i=1}^n -2(y_i - ax_i - b) = 2nb - 2 \sum_{i=1}^n (y_i - ax_i)$$

$$\frac{\partial f}{\partial b} = 0 \Rightarrow 2nb = 2 \sum_{i=1}^n (y_i - ax_i) \Rightarrow b = \frac{1}{n} \sum_{i=1}^n y_i - \frac{1}{n} a \sum_{i=1}^n x_i$$

$$\Rightarrow b = \bar{y} - a\bar{x}$$

$$\text{Similarly, } \frac{\partial f}{\partial a} = \sum_{i=1}^n -2x_i(y_i - ax_i - b) = (-2) \sum_{i=1}^n x_i y_i - (-2) \sum_{i=1}^n ax_i^2 - (-2) \sum_{i=1}^n bx_i$$

$$\frac{\partial f}{\partial a} = 0 \Rightarrow \sum_{i=1}^n x_i y_i - a \sum_{i=1}^n x_i^2 - b \sum_{i=1}^n x_i = 0$$

put  $b = \bar{y} - a\bar{x}$  into the above equation, we could get:

$$\Rightarrow \sum_{i=1}^n x_i y_i - a \sum_{i=1}^n x_i^2 - (\bar{y} - a\bar{x}) \sum_{i=1}^n x_i = 0$$

$$\Rightarrow a = \frac{\sum_{i=1}^n x_i y_i - n\bar{x}\bar{y}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2}$$

$$\text{And put } a = \frac{\sum_{i=1}^n x_i y_i - n\bar{x}\bar{y}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2} \text{ into } b = \bar{y} - a\bar{x}, \text{ we could get:}$$

$$b = \bar{y} - \frac{\sum_{i=1}^n x_i y_i - n\bar{x}\bar{y}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2} \bar{x} = \frac{\sum_{i=1}^n x_i^2 \bar{y} - \sum_{i=1}^n x_i y_i \bar{x}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2}$$

So above all, the optimal solution of  $a$  and  $b$  is:

$$a = \frac{\sum_{i=1}^n x_i y_i - n\bar{x}\bar{y}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2}$$

$$b = \bar{y} - a\bar{x} = \frac{\sum_{i=1}^n x_i^2 \bar{y} - \sum_{i=1}^n x_i y_i \bar{x}}{\sum_{i=1}^n x_i^2 - n(\bar{x})^2}$$

(c) From the analysis in (b), we could get that:

$$\frac{\partial f}{\partial b} = 2nb - 2 \sum_{i=1}^n (y_i - ax_i) = 0 \Rightarrow b = \bar{y} - a\bar{x}$$

i.e.  $b = \bar{y} - a\bar{x}$ .

Put  $(\bar{x}, \bar{y})$  into the linear model  $f(X) = aX + b$ , we could get:

$$f(\bar{x}) = a\bar{x} + b = a\bar{x} + \bar{y} - a\bar{x} = \bar{y}$$

So above all, the learned linear model  $f(X) = aX + b$  always passes through the point  $(\bar{x}, \bar{y})$ .

3. [10 points] [Regression and Classification]

- (a) When we talk about linear regression, what does 'linear' regard to? [2 points]
- (b) Assume that there are  $n$  given training examples  $\{(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)\}$ , where each input data point  $x_i$  has  $m$  real valued features. When  $m > n$ , the linear regression model is equivalent to solving an under-determined system of linear equations  $\mathbf{y} = \mathbf{X}\beta$ . One popular way to estimate  $\beta$  is to consider the so-called ridge regression:

$$\underset{\beta}{\operatorname{argmin}} \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2$$

for some  $\lambda > 0$ . This is also known as Tikhonov regularization.

Show that the optimal solution  $\beta_*$  to the above optimization problem is given by

$$\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

Hint: You need to prove that given  $\lambda > 0$ ,  $\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}$  is invertible. [5 points]

- (c) Is the given data set linear separable? If yes, construct a linear hypothesis function to separate the given data set. If no, explain the reason. [3 points]

Data	(1,3)	(4,4)	(3,-6)	(-2,1)	(-3,5)	(-6,-4)
Label	+1	-1	-1	+1	-1	-1

- (a) Linear is to the for all parameters of the regression variable  $\beta$ .

- (b) As we have learned in linear algebra, we know that the matrix  $X^T X$  must be similar and diagonalizable. i.e. there must exist a matrix  $P$  and a diagonal matrix  $\Lambda$  such that  $X^T X = P \Lambda P^{-1}$ .

Also  $\forall x \in \mathbb{R}^n$ , we have  $x^T (\mathbf{X}^T \mathbf{X}) x = (\mathbf{X}x)^T (\mathbf{X}x) = \|\mathbf{X}x\|_2^2 \geq 0$ .

So  $X^T X$  is positive semi-definite.

So all eigenvalues of  $X^T X$  are non-negative.

i.e. the diagonal matrix  $\Lambda$ 's elements are all positive.

And since  $\lambda > 0$ , so  $\lambda I$ 's all elements are all also non-negative, and  $\lambda I$  is also a diagonal matrix.

So  $X^T X + \lambda I = P \Lambda P^{-1} + \lambda P I P^{-1} = P(\Lambda + \lambda I) P^{-1}$ .

Since  $\Lambda, \lambda I$  are all diagonal matrix, so  $\Lambda + \lambda I$  is also a diagonal matrix.

And all elements in  $\Lambda + \lambda I$  are all positive, this is because in  $\Lambda$ , elements are non-negative, in  $\lambda I$ , all elements are positive. So  $\Lambda + \lambda I$  is positive defined.

Since  $X^T X + \lambda I = P(\Lambda + \lambda I) P^{-1}$ , from the knowledge of similarity and diagonalizable, we could know that  $X^T X + \lambda I$  is also positive defined.

So  $X^T X + \lambda I$  is invertible.

And let  $f(\beta) = \|\mathbf{y} - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2 = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta) + \lambda \beta^T \beta = \mathbf{y}^T \mathbf{y} - \beta^T \mathbf{X}^T \mathbf{y} - \mathbf{y}^T \mathbf{X} \beta + \beta^T \mathbf{X}^T \mathbf{X} \beta + \lambda \beta^T \beta$

Since  $f(\beta)$  is convex, so we just need to set the derivative of  $f(\beta)$  to 0 to get the optimal solution.

$$\frac{\partial f(\beta)}{\partial \beta} = 2(-\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \beta + \lambda \beta)$$

$$\frac{\partial f(\beta)}{\partial \beta} = 0 \Rightarrow (\mathbf{X}^T \mathbf{X} + \lambda I) \beta = \mathbf{X}^T \mathbf{y} \Rightarrow \beta = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$$

Since we have proved that  $X^T X + \lambda I$  is invertible, so  $(\mathbf{X}^T \mathbf{X} + \lambda I)^{-1}$  exists. So  $\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$

So above all, the optimal solution  $\beta_* = (\mathbf{X}^T \mathbf{X} + \lambda I)^{-1} \mathbf{X}^T \mathbf{y}$ .

- (c) Let the data be formed as  $\mathbf{X}_i = (x_1, x_2)$ .

And let the hypothesis function be  $f(\mathbf{X}) = b_1 x_1^2 + b_2 x_2^2 + b_3$ .

Let  $b_1 = 1, b_2 = 1, b_3 = -25$ , so the regression function is  $f(\mathbf{X}) = x_1^2 + x_2^2 - 25$ , and it's a linear regression.

Make the separate line be  $f(\mathbf{X}) = 0$ , so the separate line is  $x_1^2 + x_2^2 - 25 = 0$ .

If  $f(\mathbf{X}) \leq 0$ , set the label to be +1, else set the label to be -1.

And we could get the result as below:

Data $\mathbf{X}$	(1,3)	(4,4)	(3,-6)	(-2,1)	(-3,5)	(-6,-4)
Function value $f(\mathbf{X})$	-15	7	20	-20	9	27
Label	+1	-1	-1	+1	-1	-1

So above all, we could find that we can construct a linear hypothesis function to separate the given data set.