

HW5-Coding

January 3, 2024

1 Homework 5: Convolutional neural network (30 points)

In this part, you need to implement and train a convolutional neural network on the CIFAR-10 dataset with PyTorch. ### What is PyTorch?

PyTorch is a system for executing dynamic computational graphs over Tensor objects that behave similarly as numpy ndarray. It comes with a powerful automatic differentiation engine that removes the need for manual back-propagation.

1.0.1 Why?

- Our code will now run on GPUs! Much faster training. When using a framework like PyTorch or TensorFlow you can harness the power of the GPU for your own custom neural network architectures without having to write CUDA code directly (which is beyond the scope of this class).
- We want you to be ready to use one of these frameworks for your project so you can experiment more efficiently than if you were writing every feature you want to use by hand.
- We want you to stand on the shoulders of giants! TensorFlow and PyTorch are both excellent frameworks that will make your lives a lot easier, and now that you understand their guts, you are free to use them :)
- We want you to be exposed to the sort of deep learning code you might run into in academia or industry. ## How can I learn PyTorch?

Justin Johnson has made an excellent [tutorial](#) for PyTorch.

You can also find the detailed [API doc](#) here. If you have other questions that are not addressed by the API docs, the [PyTorch forum](#) is a much better place to ask than StackOverflow.

Install PyTorch and Skorch.

```
[ ]: !pip install -q torch skorch torchvision torchtext
```

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import skorch
```

```
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

1.1 0. Tensor Operations (5 points)

Tensor operations are important in deep learning models. In this part, you are required to get familiar to some common tensor operations in PyTorch.

1.1.1 1) Tensor squeezing, unsqueezing and viewing

Tensor squeezing, unsqueezing and viewing are important methods to change the dimension of a Tensor, and the corresponding functions are [torch.squeeze](#), [torch.unsqueeze](#) and [torch.Tensor.view](#). Please read the documents of the functions, and finish the following practice.

```
[ ]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2],
                  [3, 4],
                  [5, 6]])

x.shape

# Add two new dimensions to x by using the function torch.unsqueeze, so that
→ the size of x becomes (3, 1, 2, 1).
x = torch.unsqueeze(x, 1)
x = torch.unsqueeze(x, 3)
print(x.shape)

# Remove the two dimensions just added by using the function torch.squeeze,
→ and change the size of x back to (3, 2).
x = torch.squeeze(x, 3)
x = torch.squeeze(x, 1)
# x = torch.squeeze(x, (0, 2)) # in torch 2.0 this will work
print(x.shape)

# x is now a two-dimensional tensor, or in other words a matrix. Now use the
→ function torch.Tensor.view and change x to a one-dimensional vector with
→ size being (6).
x = x.view(6)
print(x.shape)
```

```
torch.Size([3, 1, 2, 1])
torch.Size([3, 2])
torch.Size([6])
```

1.1.2 2) Tensor concatenation and stack

Tensor concatenation and stack are operations to combine small tensors into big tensors. The corresponding functions are `torch.cat` and `torch.stack`. Please read the documents of the functions, and finish the following practice.

```
[ ]: # x is a tensor with size being (3, 2)
x = torch.Tensor([[1, 2], [3, 4], [5, 6]])

# y is a tensor with size being (3, 2)
y = torch.Tensor([[-1, -2], [-3, -4], [-5, -6]])

# Our goal is to generate a tensor z with size as (2, 3, 2), and z[0,:,:] = x,
↪ z[1,:,:] = y.

# Use torch.stack to generate such a z
z = torch.stack((x, y))
# print(z)
print(z[0,:,:])

# Use torch.cat and torch.unsqueeze to generate such a z
# print('torch.unsqueeze(x, 0) = ', torch.unsqueeze(x, 0))
z = torch.cat((torch.unsqueeze(x, 0), torch.unsqueeze(y, 0)))
# print(z)
print(z[1,:,:])

tensor([[1., 2.],
        [3., 4.],
        [5., 6.]])
tensor([[-1., -2.],
        [-3., -4.],
        [-5., -6.]])
```

1.1.3 3) Tensor expansion

Tensor expansion is to expand a tensor into a larger tensor along singleton dimensions. The corresponding functions are `torch.Tensor.expand` and `torch.Tensor.expand_as`. Please read the documents of the functions, and finish the following practice.

```
[ ]: # x is a tensor with size being (3)
x = torch.Tensor([1, 2, 3])

# Our goal is to generate a tensor z with size (2, 3), so that z[0,:,:] = x,
↪ z[1,:,:] = x.

# [TO DO]
# Change the size of x into (1, 3) by using torch.unsqueeze.
```

```

x = torch.unsqueeze(x, 0)
print(x.shape)

# [TO DO]
# Then expand the new tensor to the target tensor by using torch.Tensor.expand.
z = x.expand(2, -1)
# print(z)
print(z.shape)

```

```

torch.Size([1, 3])
torch.Size([2, 3])

```

1.1.4 4) Tensor reduction in a given dimension

In deep learning, we often need to compute the mean/sum/max/min value in a given dimension of a tensor. Please read the document of [torch.mean](#), [torch.sum](#), [torch.max](#), [torch.min](#), [torch.topk](#), and finish the following practice.

```

[ ]: # x is a random tensor with size being (10, 50)
x = torch.randn(10, 50)

# Compute the mean value for each row of x.
# You need to generate a tensor x_mean of size (10), and x_mean[k, :] is the
  ↳ mean value of the k-th row of x.

# dim = 1: eliminate the second(1)'s dimension

x_mean = torch.mean(x, dim=1)
# print(x_mean)
print(x_mean[3, ])

# Compute the sum value for each row of x.
# You need to generate a tensor x_sum of size (10).
x_sum = torch.sum(x, dim=1)
print(x_sum.shape)

# Compute the max value for each row of x.
# You need to generate a tensor x_max of size (10).
(x_max, indices) = torch.max(x, dim=1)
# print(x_max, indices)
print(x_max.shape)

# Compute the min value for each row of x.
# You need to generate a tensor x_min of size (10).
(x_min, indices) = torch.min(x, dim=1)
print(x_min.shape)

```

```

# Compute the top-5 values for each row of x.
# (wrong) You need to generate a tensor x_mean of size (10, 5), and x_top[k, :]
→ is the top-5 values of each row in x.
# (right) You need to generate a tensor, top-5 values of each row
(x_xtop, indices) = torch.topk(x, k=5, dim=1)
print((x_xtop.shape))

```

```

tensor(0.2702)
torch.Size([10])
torch.Size([10])
torch.Size([10])
torch.Size([10, 5])

```

1.2 Convolutional Neural Networks

Implement a convolutional neural network for image classification on CIFAR-10 dataset.

CIFAR-10 is an image dataset of 10 categories. Each image has a size of 32x32 pixels. The following code will download the dataset, and split it into `train` and `test`. For this question, we use the default validation split generated by Skorch.

```

[ ]: train = torchvision.datasets.CIFAR10("./data", train=True, download=True)
test = torchvision.datasets.CIFAR10("./data", train=False, download=True)

```

```

Files already downloaded and verified
Files already downloaded and verified

```

The following code visualizes some samples in the dataset. You may use it to debug your model if necessary.

```

[ ]: def plot(data, labels=None, num_sample=5):
    n = min(len(data), num_sample)
    for i in range(n):
        plt.subplot(1, n, i+1)
        plt.imshow(data[i], cmap="gray")
        plt.xticks([])
        plt.yticks([])
        if labels is not None:
            plt.title(labels[i])

train.labels = [train.classes[target] for target in train.targets]
plot(train.data, train.labels)

```



1.2.1 1) Basic CNN implementation

Consider a basic CNN model

- It has 3 convolutional layers, followed by a linear layer.
- Each convolutional layer has a kernel size of 3, a padding of 1.
- ReLU activation is applied on every hidden layer.

Please implement this model in the following section. The hyperparameters is then be tuned and you need to fill the results in the table.

a) Implement convolutional layers (10 Points) Implement the initialization function and the forward function of the CNN.

```
[ ]: class CNN(nn.Module):
    def __init__(self, channels):
        super(CNN, self).__init__()
        # implement parameter definitions here
        # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

        # print('channels = ', channels)
        self.channels = channels

        self.backbone = nn.Sequential(
            nn.Conv2d(3, channels, kernel_size=3, padding=1),
            nn.ReLU(),

            nn.Conv2d(channels, channels, kernel_size=3, padding=1),
            nn.ReLU(),

            nn.Conv2d(channels, channels, kernel_size=3, padding=1),
            nn.ReLU(),
        )

        # regard the image has the same size 32 * 32
        self.f1 = nn.Linear(channels * 32 * 32, 10)
```

```

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

def forward(self, images):
    # implement the forward function here
    # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

    images = self.backbone(images)

    # flatten the image
    images = images.view(images.shape[0], -1)
    images = self.f1(images)

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return images

```

b) Tune hyperparameters Train the CNN model on CIFAR-10 dataset. We can tune the number of channels, optimizer, learning rate and the number of epochs for best validation accuracy.

```

[ ]: # implement hyperparameters, you can select and modify the hyperparameters by
    ↪yourself here.

optimizer = [torch.optim.SGD, torch.optim.Adam]
learning_rate = [1e-3, 1e-2]
channel = [128, 256, 512]

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for o in optimizer:
        for c in channel:
            print(f'The channel was {c}, the learning rate was {l} and the optimizer
            ↪was {str(o)}')

            cnn = CNN(channels = c)

            model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
            ↪CrossEntropyLoss,

                                                device="cuda",
                                                optimizer=o,
                                                # optimizer__momentum=0.90,
                                                lr=1,
                                                max_epochs=50,
                                                batch_size=32,
                                                callbacks=[skorch.callbacks.
            ↪EarlyStopping(lower_is_better=True)])

```

```
# implement input normalization & type cast here
model.fit(train_data_normalized, np.asarray(train.targets))
```

The channel was 128, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
1	2.1158	0.3160	1.9293	
8.0129				
2	1.8676	0.3645	1.8060	
7.0036				
3	1.7704	0.3984	1.7152	
7.0300				
4	1.6917	0.4253	1.6462	
7.0520				
5	1.6336	0.4405	1.5969	
7.0703				
6	1.5857	0.4537	1.5550	
7.0006				
7	1.5434	0.4653	1.5178	
7.0303				
8	1.5062	0.4777	1.4855	
6.9867				
9	1.4740	0.4851	1.4584	
7.0018				
10	1.4458	0.4927	1.4353	
7.0111				
11	1.4201	0.4999	1.4149	
7.0238				
12	1.3961	0.5051	1.3962	
7.0131				
13	1.3729	0.5109	1.3787	
7.0701				
14	1.3496	0.5166	1.3605	
7.0898				
15	1.3254	0.5228	1.3404	
7.0641				
16	1.3003	0.5288	1.3189	
7.0448				
17	1.2751	0.5329	1.2981	
7.0350				
18	1.2510	0.5411	1.2792	
7.0506				
19	1.2286	0.5472	1.2625	
7.0371				
20	1.2082	0.5523	1.2481	
7.0452				

21	1.1893	0.5558	1.2354
7.0610			
22	1.1714	0.5615	1.2241
7.0859			
23	1.1543	0.5647	1.2139
7.0981			
24	1.1376	0.5670	1.2042
7.1177			
25	1.1213	0.5722	1.1951
7.1225			
26	1.1052	0.5741	1.1863
7.1420			
27	1.0892	0.5781	1.1781
7.1389			
28	1.0733	0.5832	1.1701
7.1626			
29	1.0574	0.5866	1.1624
7.1693			
30	1.0415	0.5895	1.1553
7.1629			
31	1.0258	0.5915	1.1485
7.1485			
32	1.0102	0.5954	1.1422
7.1612			
33	0.9948	0.5981	1.1362
7.1532			
34	0.9796	0.6016	1.1305
7.1491			
35	0.9645	0.6052	1.1253
7.1466			
36	0.9495	0.6078	1.1203
7.1506			
37	0.9347	0.6089	1.1159
7.1595			
38	0.9200	0.6109	1.1117
7.1716			
39	0.9054	0.6129	1.1079
7.1662			
40	0.8907	0.6141	1.1043
7.1629			
41	0.8761	0.6160	1.1011
7.1744			
42	0.8615	0.6161	1.0983
7.1671			
43	0.8469	0.6193	1.0959
7.1680			
44	0.8323	0.6205	1.0939
7.1720			

45	0.8176	0.6211	1.0921	
7.1611				
46	0.8030	0.6235	1.0907	
7.1709				
47	0.7884	0.6242	1.0898	
7.1742				
48	0.7738	0.6250	1.0895	
7.1943				
49	0.7593	0.6255	1.0895	7.1965
50	0.7448	0.6266	1.0902	7.1881

The channel was 256, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.0442	0.3434	1.8783	
16.4868				
2	1.8054	0.3976	1.7304	
16.5860				
3	1.7004	0.4273	1.6449	
16.6578				
4	1.6279	0.4455	1.5829	
16.6932				
5	1.5640	0.4653	1.5236	
16.7557				
6	1.5044	0.4785	1.4709	
16.7104				
7	1.4544	0.4940	1.4321	
16.7861				
8	1.4143	0.5052	1.4030	
16.7697				
9	1.3789	0.5118	1.3771	
16.7234				
10	1.3442	0.5204	1.3506	
16.7258				
11	1.3091	0.5285	1.3241	
16.5999				
12	1.2755	0.5356	1.3007	
16.5341				
13	1.2457	0.5426	1.2806	
16.4736				
14	1.2194	0.5506	1.2631	
16.4122				
15	1.1957	0.5565	1.2472	
16.3709				
16	1.1736	0.5612	1.2330	
16.4181				
17	1.1525	0.5658	1.2197	
16.4600				

18	1.1320	0.5706	1.2072	
16.4435				
19	1.1117	0.5745	1.1960	
16.4278				
20	1.0917	0.5785	1.1856	
16.4144				
21	1.0720	0.5834	1.1760	
16.4010				
22	1.0525	0.5882	1.1671	
16.3734				
23	1.0335	0.5924	1.1589	
16.3465				
24	1.0148	0.5960	1.1515	
16.3348				
25	0.9964	0.5974	1.1447	
16.3804				
26	0.9784	0.5993	1.1385	
16.3272				
27	0.9607	0.6022	1.1328	
16.3608				
28	0.9431	0.6049	1.1276	
16.3947				
29	0.9255	0.6067	1.1225	
16.4668				
30	0.9080	0.6086	1.1179	
16.5375				
31	0.8904	0.6086	1.1137	16.5316
32	0.8726	0.6107	1.1097	
16.6019				
33	0.8547	0.6129	1.1059	
16.6545				
34	0.8366	0.6152	1.1026	
16.6232				
35	0.8184	0.6159	1.0996	
16.6570				
36	0.8000	0.6172	1.0973	
16.5988				
37	0.7814	0.6189	1.0957	
16.5260				
38	0.7628	0.6214	1.0950	
16.5240				
39	0.7441	0.6227	1.0952	16.4719
40	0.7254	0.6224	1.0965	16.4138
41	0.7067	0.6222	1.0988	16.4675
42	0.6880	0.6223	1.1025	16.4783

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.001 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.9867	0.3623	1.8191	
46.1175				
2	1.7513	0.4171	1.6693	
46.1142				
3	1.6400	0.4468	1.5838	
45.7091				
4	1.5603	0.4680	1.5144	
45.3723				
5	1.4935	0.4853	1.4582	
45.4629				
6	1.4396	0.5000	1.4170	
45.4543				
7	1.3952	0.5077	1.3844	
45.3653				
8	1.3541	0.5188	1.3529	
45.4058				
9	1.3127	0.5287	1.3208	
45.8955				
10	1.2729	0.5375	1.2915	
46.3144				
11	1.2376	0.5454	1.2668	
46.4242				
12	1.2072	0.5545	1.2464	
46.3143				
13	1.1800	0.5601	1.2289	
45.7325				
14	1.1545	0.5658	1.2133	
45.9819				
15	1.1300	0.5724	1.1989	
45.9304				
16	1.1059	0.5773	1.1857	
45.4933				
17	1.0821	0.5817	1.1732	
45.3715				
18	1.0587	0.5876	1.1618	
45.4426				
19	1.0358	0.5915	1.1513	
45.5602				
20	1.0134	0.5973	1.1418	
46.1363				
21	0.9915	0.6005	1.1330	
46.3644				
22	0.9700	0.6039	1.1250	
45.7330				
23	0.9487	0.6072	1.1174	
45.8170				

24	0.9277	0.6107	1.1103	
45.4220				
25	0.9067	0.6116	1.1037	
45.4124				
26	0.8858	0.6138	1.0975	
45.3856				
27	0.8651	0.6175	1.0919	
45.7772				
28	0.8444	0.6187	1.0870	
46.3898				
29	0.8239	0.6195	1.0828	
46.3944				
30	0.8035	0.6216	1.0797	
46.1950				
31	0.7833	0.6235	1.0776	
45.4395				
32	0.7632	0.6250	1.0765	
45.5423				
33	0.7433	0.6255	1.0766	45.4350
34	0.7235	0.6264	1.0779	45.3818
35	0.7037	0.6283	1.0803	45.3848
36	0.6839	0.6295	1.0839	45.7210

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.4839	0.5829	1.1889	
7.5850				
2	1.0473	0.6502	1.0072	
7.5615				
3	0.8172	0.6484	1.0422	7.5612
4	0.6176	0.6250	1.2363	7.5515
5	0.4350	0.5978	1.5414	7.5729
6	0.3217	0.6011	1.8818	7.5781

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	1.7182	0.5077	1.3801	
16.8513				
2	1.2644	0.5565	1.2567	
16.8689				
3	1.0805	0.5825	1.1950	
16.9014				
4	0.8733	0.5819	1.2774	16.8458
5	0.6655	0.5767	1.5078	16.8153

6	0.5032	0.5598	1.8750	16.8164
7	0.3858	0.5273	2.4469	16.8811

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 512, the learning rate was 0.001 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.3566	0.1000	2.3026	
46.3657				
2	2.3028	0.1000	2.3026	46.2071
3	2.3028	0.1000	2.3026	45.8301
4	2.3028	0.1000	2.3026	46.2545
5	2.3028	0.1000	2.3026	46.7461

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 128, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7802	0.4473	1.5509	
7.2227				
2	1.4604	0.4953	1.3926	
7.1891				
3	1.3259	0.5312	1.2938	
7.1699				
4	1.2092	0.5520	1.2376	
7.1683				
5	1.1174	0.5713	1.1887	
7.1695				
6	1.0313	0.5887	1.1468	
7.1558				
7	0.9441	0.6095	1.1049	
7.1596				
8	0.8574	0.6212	1.0799	
7.1328				
9	0.7744	0.6294	1.0816	7.1242
10	0.6936	0.6324	1.1068	7.1195
11	0.6121	0.6323	1.1574	7.1237
12	0.5278	0.6290	1.2431	7.1241

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 256, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7304	0.4675	1.4906	
16.3829				
2	1.3863	0.5328	1.3044	
16.5476				
3	1.2254	0.5659	1.2175	

16.5300				
4	1.1000	0.5942	1.1371	
16.5410				
5	0.9843	0.6187	1.0781	
16.5368				
6	0.8811	0.6341	1.0504	
16.4840				
7	0.7862	0.6431	1.0436	
16.4648				
8	0.6937	0.6450	1.0633	16.3988
9	0.5998	0.6420	1.1201	16.3038
10	0.5008	0.6338	1.2165	16.3112
11	0.3968	0.6297	1.3476	16.3209

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.01 and the optimizer was <class 'torch.optim.sgd.SGD'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.7122	0.4814	1.4455	
45.7170				
2	1.3445	0.5414	1.2749	
45.8045				
3	1.1590	0.5824	1.1645	
46.0595				
4	1.0109	0.6085	1.0983	
46.3534				
5	0.8835	0.6283	1.0726	
46.3809				
6	0.7679	0.6391	1.0765	46.4801
7	0.6554	0.6344	1.1211	46.4627
8	0.5380	0.6302	1.2184	45.9412
9	0.4127	0.6206	1.3871	45.4630

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.9942	0.1000	2.3039	
7.4479				
2	2.3042	0.1000	2.3039	7.4428
3	2.3042	0.1000	2.3039	7.4505
4	2.3042	0.1000	2.3039	7.4814
5	2.3042	0.1000	2.3039	7.4697

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----

1	8.8825	0.1000	2.3039	
16.5144				
2	2.3042	0.1000	2.3039	16.7190
3	2.3042	0.1000	2.3039	16.8494
4	2.3042	0.1000	2.3039	16.9722
5	2.3042	0.1000	2.3039	17.0036

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.01 and the optimizer was <class 'torch.optim.adam.Adam'>

epoch	train_loss	valid_acc	valid_loss	dur
1	55.4451	0.1000	2.3038	
46.9036				
2	2.3042	0.1000	2.3039	46.8506
3	2.3042	0.1000	2.3039	46.5957
4	2.3042	0.1000	2.3039	45.6649
5	2.3042	0.1000	2.3039	45.5824

Stopping since valid_loss has not improved in the last 5 epochs.

Write down **validation accuracy** of your model under different hyperparameter settings. Note the validation set is automatically split by Skorch during `model.fit()`.

learning rate = 10^{-3} : | #channel for each layer | optimizer | SGD | Adam | |:-----|
 -----:|:-----:|:-----:| | 128 |0.6266|0.6502| | 256 |0.6227|0.5825| | 512 |0.6295|0.1000|

learning rate = 10^{-2} : | #channel for each layer | optimizer | SGD | Adam | |:-----|
 -----:|:-----:|:-----:| | 128 |0.6324|0.1000| | 256 |0.6450|0.1000| | 512 |0.6391|0.1000|

1.2.2 2) Full CNN implementation (10 points)

Based on the CNN in the previous question, implement a full CNN model with max pooling layer.

- Add a max pooling layer after each convolutional layer.
- Each max pooling layer has a kernel size of 2 and a stride of 2.

Please implement this model in the following section. The hyperparameters is then be tuned and fill the results in the table. You are also required to complete the questions.

a) Implement max pooling layers Similar to the CNN implementation in previous question, implement max pooling layers.

```
[ ]: class CNN_MaxPool(nn.Module):
      def __init__(self, channels):
          super(CNN_MaxPool, self).__init__()
          # implement parameter definitions here
          # *****START OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

          self.channels = channels
```



```

self.backbone = nn.Sequential(
    nn.Conv2d(3, channels, kernel_size=3, padding=1),
    nn.BatchNorm2d(channels),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Dropout2d(p=0.2),

    nn.Conv2d(channels, channels, kernel_size=3, padding=1),
    nn.BatchNorm2d(channels),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Dropout2d(p=0.2),

    nn.Conv2d(channels, channels, kernel_size=3, padding=1),
    nn.BatchNorm2d(channels),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2),
    nn.Dropout2d(p=0.2)
)

# regard the image has the same size 32 * 32
# after 3 max pooling, the size is 32 / 2 / 2 / 2 = 4
self.f1 = nn.Linear(channels * 4 * 4, 10)

# *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****

def forward(self, images):
    # implement the forward function here

    images = self.backbone(images)

    # flatten the image
    images = images.view(images.shape[0], -1)

    images = self.f1(images)

    # *****END OF YOUR CODE (DO NOT DELETE/MODIFY THIS LINE)*****
    return images

```

b) Tune hyperparameters Based on the better optimizer found in the previous problem, we can tune the number of channels and learning rate for best validation accuracy.

```

[ ]: # implement hyperparameters, you can select and modify the hyperparameters by
      ↪yourself here.
learning_rate = [1e-4, 1e-3, 1e-2]
channel = [128, 256, 512]

```

*# Select the better optimizer by the result shown in the previous problem, you
 ↳ can select and modify it by yourself here.*

```

better_optimizer = torch.optim.Adam

train_data_normalized = torch.Tensor(train.data/255)
train_data_normalized = train_data_normalized.permute(0,3,1,2)

for l in learning_rate:
    for c in channel:
        print(f'The channel was {c}, the learning rate was {l}')

        cnn = CNN_MaxPool(channels = c)

        model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.
↳CrossEntropyLoss,
                                         device="cuda",
                                         optimizer=better_optimizer,
                                         lr=l,
                                         max_epochs=50,
                                         batch_size=32 * 8,
                                         callbacks=[skorch.callbacks.
↳EarlyStopping(lower_is_better=True)])
        # implement input normalization & type cast here
        model.fit(train_data_normalized, np.asarray(train.targets))
  
```

The channel was 128, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.8045	0.4908	1.4717	
3.2317				
2	1.4935	0.5442	1.3015	
3.1580				
3	1.3585	0.5886	1.1966	
3.1443				
4	1.2596	0.6146	1.1193	
3.1602				
5	1.1847	0.6342	1.0625	
3.1622				
6	1.1234	0.6518	1.0135	
3.1547				
7	1.0764	0.6625	0.9778	
3.1650				
8	1.0358	0.6719	0.9484	
3.1638				
9	0.9967	0.6813	0.9210	
3.1614				

10	0.9698	0.6923	0.8992	
3.1662				
11	0.9418	0.6989	0.8795	
3.1710				
12	0.9159	0.7008	0.8661	
3.1699				
13	0.8987	0.7088	0.8503	
3.1751				
14	0.8736	0.7155	0.8327	
3.1795				
15	0.8550	0.7192	0.8209	
3.1766				
16	0.8351	0.7245	0.8086	
3.1782				
17	0.8160	0.7282	0.7994	
3.1833				
18	0.8006	0.7312	0.7842	
3.1876				
19	0.7876	0.7328	0.7803	
3.1829				
20	0.7678	0.7339	0.7716	
3.1953				
21	0.7583	0.7398	0.7576	
3.1882				
22	0.7469	0.7415	0.7519	
3.1923				
23	0.7304	0.7401	0.7504	3.1926
24	0.7186	0.7418	0.7470	
3.1946				
25	0.7094	0.7458	0.7346	
3.1994				
26	0.6931	0.7453	0.7383	3.1950
27	0.6847	0.7467	0.7313	
3.1931				
28	0.6752	0.7515	0.7170	
3.1928				
29	0.6627	0.7544	0.7119	
3.1875				
30	0.6524	0.7511	0.7180	3.1824
31	0.6418	0.7569	0.7052	
3.1839				
32	0.6332	0.7589	0.6995	
3.1801				
33	0.6226	0.7579	0.6945	3.1792
34	0.6115	0.7575	0.6946	3.1858
35	0.6090	0.7609	0.6869	
3.1795				
36	0.5965	0.7622	0.6868	

3.1880				
37	0.5886	0.7613	0.6831	3.1852
38	0.5806	0.7669	0.6744	
3.1756				
39	0.5691	0.7656	0.6816	3.1742
40	0.5664	0.7697	0.6742	
3.1783				
41	0.5545	0.7657	0.6778	3.1825
42	0.5465	0.7708	0.6708	
3.1714				
43	0.5400	0.7692	0.6658	3.1728
44	0.5343	0.7663	0.6744	3.1602
45	0.5261	0.7737	0.6608	
3.1719				
46	0.5252	0.7717	0.6618	3.1673
47	0.5109	0.7700	0.6638	3.1695
48	0.5050	0.7716	0.6618	3.1599
49	0.4969	0.7755	0.6571	
3.1588				
50	0.4920	0.7733	0.6619	3.1607
The channel was 256, the learning rate was 0.0001				
epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6336	0.5527	1.2881	
6.1969				
2	1.2999	0.6045	1.1220	
6.1984				
3	1.1574	0.6433	1.0264	
6.2124				
4	1.0544	0.6692	0.9555	
6.2015				
5	0.9872	0.6815	0.9150	
6.2078				
6	0.9240	0.6962	0.8753	
6.1982				
7	0.8772	0.7051	0.8505	
6.2226				
8	0.8395	0.7134	0.8215	
6.1924				
9	0.8040	0.7197	0.7986	
6.1769				
10	0.7724	0.7305	0.7737	
6.1587				
11	0.7434	0.7366	0.7681	
6.1682				
12	0.7170	0.7445	0.7461	
6.1463				
13	0.6899	0.7460	0.7380	

6.1383				
14	0.6647	0.7519	0.7267	
6.1318				
15	0.6392	0.7543	0.7104	
6.1295				
16	0.6187	0.7574	0.7044	
6.1260				
17	0.5990	0.7591	0.6985	
6.1007				
18	0.5774	0.7591	0.6911	6.0768
19	0.5607	0.7651	0.6901	
6.0741				
20	0.5382	0.7668	0.6811	
6.1156				
21	0.5218	0.7682	0.6737	
6.1021				
22	0.5011	0.7710	0.6630	
6.1216				
23	0.4857	0.7717	0.6631	6.0929
24	0.4669	0.7751	0.6545	
6.1075				
25	0.4572	0.7752	0.6540	
6.1212				
26	0.4396	0.7743	0.6600	6.1145
27	0.4239	0.7756	0.6529	
6.1094				
28	0.4048	0.7788	0.6474	
6.1126				
29	0.3948	0.7783	0.6507	6.1339
30	0.3807	0.7789	0.6526	6.1389
31	0.3655	0.7767	0.6542	6.1480
32	0.3515	0.7814	0.6444	
6.1751				
33	0.3393	0.7840	0.6455	6.1548
34	0.3247	0.7816	0.6497	6.1603
35	0.3153	0.7823	0.6487	6.1681
36	0.3007	0.7860	0.6388	
6.1728				
37	0.2920	0.7851	0.6449	6.1701
38	0.2808	0.7826	0.6554	6.1807
39	0.2743	0.7845	0.6524	6.1956
40	0.2562	0.7842	0.6490	6.2041

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 512, the learning rate was 0.0001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.4867	0.5875	1.1599	
13.7566				

2	1.1323	0.6544	0.9950	
13.8768				
3	0.9813	0.6778	0.9199	
13.9526				
4	0.8851	0.7036	0.8458	
14.0137				
5	0.8155	0.7077	0.8244	
14.0539				
6	0.7579	0.7197	0.7897	
14.0974				
7	0.7093	0.7321	0.7623	
14.0722				
8	0.6648	0.7351	0.7490	
14.0787				
9	0.6232	0.7519	0.7108	
14.0801				
10	0.5813	0.7579	0.6961	
13.9528				
11	0.5446	0.7622	0.6848	
13.9407				
12	0.5114	0.7663	0.6737	
14.0248				
13	0.4788	0.7650	0.6776	13.9431
14	0.4443	0.7624	0.6877	14.0306
15	0.4155	0.7715	0.6587	
14.0328				
16	0.3874	0.7740	0.6575	
13.8163				
17	0.3581	0.7736	0.6607	13.7884
18	0.3337	0.7801	0.6459	
13.7710				
19	0.3078	0.7799	0.6515	13.7418
20	0.2844	0.7740	0.6669	13.6890
21	0.2644	0.7758	0.6733	13.6630
22	0.2433	0.7715	0.6883	13.8098

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 128, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
1	1.5918	0.5432	1.2453	
3.1035				
2	1.1911	0.6467	1.0016	
3.1075				
3	1.0245	0.6690	0.9421	
3.1208				
4	0.9245	0.7108	0.8356	
3.1018				
5	0.8580	0.7191	0.8196	

3.1075				
6	0.8062	0.7294	0.7835	
3.1001				
7	0.7627	0.7342	0.7764	
3.0934				
8	0.7159	0.7441	0.7422	
3.0959				
9	0.6845	0.7535	0.7172	
3.1002				
10	0.6500	0.7504	0.7301	3.1135
11	0.6233	0.7222	0.8134	3.1064
12	0.5967	0.7500	0.7301	3.1062
13	0.5679	0.7382	0.7676	3.0977
14	0.5419	0.7573	0.7070	
3.1099				
15	0.5189	0.7275	0.8109	3.1219
16	0.5012	0.7383	0.7763	3.1253
17	0.4856	0.7759	0.6584	
3.1117				
18	0.4610	0.7671	0.6920	3.1101
19	0.4503	0.7484	0.7565	3.1313
20	0.4330	0.7603	0.7146	3.1280
21	0.4164	0.7584	0.7378	3.1179

Stopping since valid_loss has not improved in the last 5 epochs.

The channel was 256, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.6579	0.5450	1.2561	
6.1184				
2	1.1583	0.6407	1.0060	
6.1430				
3	0.9712	0.6851	0.9020	
6.1335				
4	0.8689	0.7178	0.8076	
6.1475				
5	0.7908	0.7155	0.8156	6.1657
6	0.7173	0.7254	0.7869	
6.1685				
7	0.6691	0.7162	0.8146	6.1811
8	0.6200	0.7231	0.8026	6.2037
9	0.5697	0.7593	0.6993	
6.1986				
10	0.5247	0.7511	0.7574	6.2029
11	0.4784	0.7777	0.6565	
6.2012				
12	0.4481	0.7608	0.7093	6.2085
13	0.4193	0.7678	0.7088	6.2271
14	0.3876	0.7664	0.7174	6.2404

15 0.3628 0.7653 0.7162 6.2230
 Stopping since valid_loss has not improved in the last 5 epochs.
 The channel was 512, the learning rate was 0.001

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	1.8660	0.5239	1.3267	
13.8909				
2	1.2144	0.6288	1.0492	
14.0318				
3	0.9933	0.6585	1.0033	
14.0499				
4	0.8716	0.7019	0.8538	
14.0592				
5	0.7827	0.7338	0.7637	
13.9973				
6	0.7105	0.7051	0.8631	13.9783
7	0.6455	0.7081	0.8460	14.0149
8	0.5830	0.7151	0.8477	14.0355
9	0.5239	0.7465	0.7621	
14.0760				
10	0.4676	0.7337	0.8000	14.0717
11	0.4258	0.7354	0.8330	13.8616
12	0.4056	0.7492	0.8092	13.8587
13	0.3547	0.7633	0.7739	13.9999

Stopping since valid_loss has not improved in the last 5 epochs.
 The channel was 128, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	2.4661	0.3902	1.6711	
3.1283				
2	1.4921	0.5048	1.3910	
3.1077				
3	1.2807	0.5838	1.1837	
3.1125				
4	1.1125	0.6630	0.9687	
3.0927				
5	0.9888	0.6790	0.9151	
3.0898				
6	0.9001	0.6382	1.0577	3.0995
7	0.8429	0.6774	0.9419	3.0966
8	0.7930	0.7120	0.8200	
3.0852				
9	0.7398	0.6705	0.9908	3.0919
10	0.7032	0.7199	0.8102	
3.0905				
11	0.6733	0.7302	0.7912	
3.0858				
12	0.6391	0.7149	0.8442	3.0781

13	0.6147	0.7170	0.8561	3.0809
14	0.5882	0.7029	0.8850	3.0770
15	0.5654	0.7282	0.8245	3.0779

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 256, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	3.5919	0.3983	1.6982	
6.0360				
2	1.4422	0.4942	1.4063	
6.0587				
3	1.2272	0.6151	1.1065	
6.0509				
4	1.0562	0.6300	1.0532	
6.0635				
5	0.9361	0.6477	1.0331	
6.0767				
6	0.8409	0.6235	1.1104	6.0766
7	0.7679	0.7008	0.8644	
6.0790				
8	0.6919	0.6667	0.9843	6.0795
9	0.6367	0.6870	0.9418	6.0814
10	0.5896	0.7083	0.8764	6.0957
11	0.5484	0.7123	0.8610	
6.1184				
12	0.4997	0.7515	0.7285	
6.1226				
13	0.4751	0.7501	0.7545	6.1230
14	0.4319	0.7270	0.8634	6.1406
15	0.4018	0.7519	0.7678	6.1402
16	0.3827	0.7382	0.8527	6.1434

Stopping since valid_loss has not improved in the last 5 epochs.
The channel was 512, the learning rate was 0.01

epoch	train_loss	valid_acc	valid_loss	dur
-----	-----	-----	-----	-----
1	6.0569	0.4086	1.6634	
13.7351				
2	1.5010	0.4979	1.4120	
13.8113				
3	1.2651	0.5863	1.1882	
13.7868				
4	1.0987	0.6096	1.1351	
13.8378				
5	0.9611	0.6584	0.9870	
13.8555				
6	0.8590	0.6933	0.8747	
13.9215				
7	0.7620	0.7215	0.8052	

13.9019				
8	0.6876	0.7403	0.7532	
13.9872				
9	0.6161	0.7516	0.7395	
13.9732				
10	0.5579	0.6949	0.9249	13.9310
11	0.4903	0.6445	1.1340	13.9594
12	0.4353	0.7367	0.7914	13.9665
13	0.3780	0.7526	0.7753	13.9434

Stopping since valid_loss has not improved in the last 5 epochs.

Write down the **validation accuracy** of the model under different hyperparameter settings.

learning rate = 10^{-4} : | #channel for each layer | validation accuracy | |:-----:|:-----:
 -----:| | 128 | 0.7755 | | 256 | 0.7860 | | 512 | 0.7801 |

learning rate = 10^{-3} : | #channel for each layer | validation accuracy | |:-----:|:-----:
 -----:| | 128 | 0.7759 | | 256 | 0.7777 | | 512 | 0.7633 |

learning rate = 10^{-2} : | #channel for each layer | validation accuracy | |:-----:|:-----:
 -----:| | 128 | 0.7302 | | 256 | 0.7519 | | 512 | 0.7526 |

For the best model you have, test it on the test set.

```
[ ]: # implement the same input normalization & type cast here

optimizer = torch.optim.Adam
learning_rate = 1e-4
channel = 256

cnn = CNN_MaxPool(channels = channel)
model = skorch.NeuralNetClassifier(cnn, criterion=torch.nn.CrossEntropyLoss,
                                  device="cuda",
                                  optimizer=optimizer,
                                  lr=learning_rate,
                                  max_epochs=50,
                                  batch_size=32 * 8,
                                  callbacks=[skorch.callbacks.
→EarlyStopping(lower_is_better=True)])

# implement input normalization & type cast here
model.fit(train_data_normalized, np.asarray(train.targets))

test_data_normalized = torch.Tensor(test.data/255)
test_data_normalized = test_data_normalized.permute(0,3,1,2)
test_predictions = model.predict(test_data_normalized)
sklearn.metrics.accuracy_score(test.targets, test_predictions)
```

epoch	train_loss	valid_acc	valid_loss	dur
-------	------------	-----------	------------	-----

-----	-----	-----	-----	-----
1	1.6229	0.5552	1.2808	
5.7887				
2	1.2905	0.6147	1.1130	
5.7904				
3	1.1429	0.6455	1.0174	
5.7758				
4	1.0483	0.6692	0.9528	
5.7884				
5	0.9770	0.6930	0.8981	
5.7653				
6	0.9212	0.7024	0.8618	
5.7668				
7	0.8761	0.7114	0.8418	
5.8012				
8	0.8357	0.7177	0.8171	
5.7866				
9	0.8008	0.7259	0.7928	
5.7409				
10	0.7699	0.7294	0.7760	
5.7591				
11	0.7416	0.7389	0.7623	
5.7847				
12	0.7144	0.7436	0.7479	
5.7757				
13	0.6867	0.7459	0.7379	
5.7796				
14	0.6621	0.7522	0.7261	
5.7849				
15	0.6387	0.7562	0.7117	
5.7848				
16	0.6178	0.7578	0.7046	
5.7691				
17	0.5950	0.7620	0.6940	
5.7612				
18	0.5779	0.7624	0.6915	
5.7960				
19	0.5589	0.7671	0.6826	
5.7776				
20	0.5397	0.7683	0.6823	
5.7812				
21	0.5170	0.7703	0.6769	
5.7648				
22	0.5039	0.7711	0.6676	
5.7785				
23	0.4839	0.7740	0.6576	
5.7657				
24	0.4651	0.7744	0.6569	

5.7748				
25	0.4493	0.7770	0.6514	
5.7814				
26	0.4353	0.7772	0.6568	5.7899
27	0.4183	0.7785	0.6526	5.7753
28	0.4051	0.7792	0.6506	
5.7794				
29	0.3886	0.7831	0.6463	
5.7860				
30	0.3738	0.7800	0.6461	5.7802
31	0.3578	0.7841	0.6455	
5.7607				
32	0.3497	0.7830	0.6408	5.7976
33	0.3381	0.7836	0.6397	5.7657
34	0.3251	0.7826	0.6426	5.7709
35	0.3099	0.7844	0.6464	5.7790
36	0.2998	0.7850	0.6435	5.7733
37	0.2838	0.7862	0.6367	
5.7886				
38	0.2772	0.7858	0.6401	5.7790
39	0.2692	0.7871	0.6405	5.8318
40	0.2595	0.7866	0.6365	5.7907
41	0.2484	0.7872	0.6409	5.7713
42	0.2385	0.7827	0.6520	5.7901
43	0.2300	0.7860	0.6646	5.7787
44	0.2202	0.7845	0.6562	5.7917

Stopping since valid_loss has not improved in the last 5 epochs.

[]: 0.7804

How much **test accuracy** do you get? What can you conclude for the design of CNN structure and tuning of hyperparameters? (5 points)

Your Answer:

1. The test accuracy we got is 0.7804.
The best model is the one with 256 channels for each layer and learning rate 10^{-4} .
2. From the results above, we can conclude that the number of channels, learning rate and the optimizer are important hyperparameters, they could greatly influence the training speed and evaluation effects. Also the structure of CNN is also important, the max pooling layer could improve the performance of the model. And with dropout layer, the model could be more robust to against overfitting.
The number of channels which also influence the net structure, could effect performance as hyperparameter. It should not be too small or too large. The learning rate should not be too large, otherwise the model will not converge.